

# Movie Recommendation System using MovieLens Dataset

Tim Chu and Elliot Frost

# MovieLens Dataset Info

- 100,000 ratings and 3,600 tag applications applied to 9,000 movies by 600 users
- ratings.csv
  - userId's, movielens movieId's, and ratings
- movies.csv
  - movielens movieId's, titles, and genres
- links.csv
  - movielens movieId's, imdbId's, and tmdbId's
- tags.csv
  - userId's, movielens movieId's, and user submitted tags

# The Problem

- We wanted to see if we would be able to predict if a user would like a movie
  - We extended that idea by predicting a user's star rating of a movie
- Lots of data was provided, but what would be relevant?
  - tags.csv - Using natural language processing to generalize tags would've been possible
    - Outside of the scope of our project
  - links.csv - We didn't care too much about imdb and tmdb's movie id's

# The Solution

- Collaborative Filtering!
  - Average the reviews from other likeminded people
  - Check our user's reviews from movies with similar genres

# Challenges

- Which distance function to use?
  - We tested the following:
    - Pearson Correlation
    - Cosine Similarity
    - Euclidean Distance
- What language to use?
  - We both weren't too familiar with any of the programming languages suitable for data mining
    - Python!

# Results


- Ran a cross-validation test routine to calculate the Root Mean Square Error for each distance measure

Cosine RMSE:	Pearson RMSE:	Euclidean RMSE:
0.687	0.506	0.891

- Pearson Correlation gave us the best results
  - Our default for recommender system

# Results cont.

Movie Ratings



Rating Accuracy


Pearson: 87.8%

Cosine: 85.2%

Euclidean: 83.4%

# Results cont.

Movie Ratings



Rating Accuracy

Pearson: 76.5%

Cosine: 75.5%

Euclidean: 72.2%



Thank You!